

Transfer_Learning_multiclass

May 14, 2019

1 Getting Started

This section sets up the environment for access to the Universal Sentence Encoder on TF Hub and provides examples of applying the encoder to words, sentences, and paragraphs. Tutorial followed: <https://www.dlology.com/blog/keras-meets-universal-sentence-encoder-transfer-learning-for-text-data/>

```
In [1]: # Install the latest Tensorflow version.
!pip3 install --quiet "tensorflow>=1.7"
# Install TF-Hub.
!pip3 install --quiet tensorflow-hub
!pip3 install seaborn
```

```
Requirement already satisfied: seaborn in /usr/local/lib/python3.6/dist-packages (0.9.0)
Requirement already satisfied: numpy>=1.9.3 in /usr/local/lib/python3.6/dist-packages (from se
Requirement already satisfied: pandas>=0.15.2 in /usr/local/lib/python3.6/dist-packages (from s
Requirement already satisfied: matplotlib>=1.4.3 in /usr/local/lib/python3.6/dist-packages (fr
Requirement already satisfied: scipy>=0.14.0 in /usr/local/lib/python3.6/dist-packages (from s
Requirement already satisfied: python-dateutil>=2.5.0 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: pytz>=2011k in /usr/local/lib/python3.6/dist-packages (from pan
Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.6/dist-packages (from ma
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.6/dist-packages (fr
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/pytho
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.6/dist-packages (from python
Requirement already satisfied: setuptools in /usr/local/lib/python3.6/dist-packages (from kiwi
```

More detailed information about installing Tensorflow can be found at <https://www.tensorflow.org/install/>.

```
In [2]: import tensorflow as tf
import tensorflow_hub as hub
import matplotlib.pyplot as plt
import numpy as np
import os
import pandas as pd
import re
```

```
WARNING: Logging before flag parsing goes to stderr.
W0513 08:26:49.199342 139779932305280 __init__.py:56] Some hub symbols are not available because
Using TensorFlow backend.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/control_flow_ops.py:1144: tf.nn.nn_ops.nn_max_pool (from tensorflow.python.ops.nn_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Colocations handled automatically by placer.
```

```
/usr/local/lib/python3.6/dist-packages/pandas/core/generic.py:5096: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html
    self[name] = value
```

```
In [9]: category_counts = len(df_train.label.cat.categories)
        category_counts
```

```
Out[9]: 3
```

```
In [0]: def UniversalEmbedding(x):
        return embed(tf.squeeze(tf.cast(x, tf.string)), signature="default", as_dict=True)
```

```
In [28]: input_text = layers.Input(shape=(1,), dtype=tf.string)
        embedding = layers.Lambda(UniversalEmbedding, output_shape=(embed_size,))(input_text)
        dense = layers.Dense(256, activation='relu')(embedding)
        pred = layers.Dense(category_counts, activation='softmax')(dense)
        model = Model(inputs=[input_text], outputs=pred)
        model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
        model.summary()
```

```
INFO:tensorflow:Saver not created because there are no variables in the graph to restore
```

```
I0513 08:42:20.056455 139779932305280 saver.py:1483] Saver not created because there are no variables to save
```

```
-----
Layer (type)                 Output Shape              Param #
=====
input_4 (InputLayer)         (None, 1)                 0
-----
lambda_3 (Lambda)           (None, 512)               0
-----
dense_5 (Dense)              (None, 256)              131328
-----
dense_6 (Dense)              (None, 3)                 771
=====
Total params: 132,099
Trainable params: 132,099
Non-trainable params: 0
-----
```

```
In [0]: hate=df_train[df_train.label==0]
        offensive=df_train[df_train.label==1]
        neutral=df_train[df_train.label==2]
```

```

In [0]: hate=hate.sample(frac=1)
        offensive=offensive.sample(frac=1)
        neutral=neutral.sample(frac=1)

In [0]: new_hate=hate[:min(len(hate), len(offensive), len(neutral))]
        new_off=offensive[:min(len(hate), len(offensive), len(neutral))]
        new_neutral=neutral[:min(len(hate), len(offensive), len(neutral))]

In [0]: new_hate = new_hate.append(new_off)
        new_hate=new_hate.append(new_neutral)
        df_train=new_hate

In [0]: train_text = df_train['text'].tolist()
        train_text = np.array(train_text, dtype=object)[: , np.newaxis]

        train_label = np.asarray(pd.get_dummies(df_train.label), dtype = np.int8)

In [21]: train_text.shape

Out[21]: (4290, 1)

In [23]: train_label.shape

Out[23]: (4290, 3)

```

1.1 Train Keras model and save weights

This only train and save our Keras layers not the embed module' weights.

```

In [29]: with tf.Session() as session:
        K.set_session(session)
        session.run(tf.global_variables_initializer())
        session.run(tf.tables_initializer())
        history = model.fit(train_text,
                            train_label,
                            validation_split=0.2,
                            epochs=15,
                            batch_size=32)
        model.save_weights('./model.h5')

```

Train on 3432 samples, validate on 858 samples

Epoch 1/15

3432/3432 [=====] - 8s 2ms/step - loss: 0.6982 - acc: 0.7165 - val_loss: 0.5297

Epoch 2/15

3432/3432 [=====] - 4s 1ms/step - loss: 0.5297 - acc: 0.7788 - val_loss: 0.5018

Epoch 3/15

3432/3432 [=====] - 4s 1ms/step - loss: 0.5018 - acc: 0.7952 - val_loss: 0.4871

Epoch 4/15

3432/3432 [=====] - 3s 957us/step - loss: 0.4871 - acc: 0.7966 - val_loss: 0.4871

```

Epoch 5/15
3432/3432 [=====] - 3s 927us/step - loss: 0.4701 - acc: 0.8086 - val_
Epoch 6/15
3432/3432 [=====] - 3s 936us/step - loss: 0.4565 - acc: 0.8167 - val_
Epoch 7/15
3432/3432 [=====] - 3s 934us/step - loss: 0.4396 - acc: 0.8223 - val_
Epoch 8/15
3432/3432 [=====] - 3s 940us/step - loss: 0.4292 - acc: 0.8234 - val_
Epoch 9/15
3432/3432 [=====] - 3s 949us/step - loss: 0.4122 - acc: 0.8342 - val_
Epoch 10/15
3432/3432 [=====] - 3s 938us/step - loss: 0.3945 - acc: 0.8360 - val_
Epoch 11/15
3432/3432 [=====] - 3s 988us/step - loss: 0.3769 - acc: 0.8470 - val_
Epoch 12/15
3432/3432 [=====] - 4s 1ms/step - loss: 0.3703 - acc: 0.8473 - val_lo
Epoch 13/15
3432/3432 [=====] - 4s 1ms/step - loss: 0.3438 - acc: 0.8663 - val_lo
Epoch 14/15
3432/3432 [=====] - 3s 965us/step - loss: 0.3295 - acc: 0.8677 - val_
Epoch 15/15
3432/3432 [=====] - 3s 940us/step - loss: 0.3116 - acc: 0.8811 - val_

```

1.2 Make predictions

```

In [0]: # new_text = list(s.comments)
        # new_text = np.array(new_text, dtype=object)[: , np.newaxis]
        new_text=train_text
        with tf.Session() as session:
            K.set_session(session)
            session.run(tf.global_variables_initializer())
            session.run(tf.tables_initializer())
            model.load_weights('./model.h5')
            predicts = model.predict(new_text, batch_size=32)

In [0]: categories = df_train.label.cat.categories.tolist()
        predict_logits = predicts.argmax(axis=1)
        predict_labels = [categories[logit] for logit in predict_logits]
        predict_labels

In [40]: df_train.label

Out[40]: 10929    0
          731    0
          15654   0
          6487   0
          5362   0
          14198   0

```

17757	0
23325	0
2704	0
6348	0
5513	0
18038	0
3006	0
13413	0
12254	0
2521	0
13722	0
20266	0
6454	0
6683	0
18676	0
15269	0
5580	0
11681	0
3902	0
9066	0
14591	0
2472	0
19351	0
7029	0
	..
15195	2
6077	2
18846	2
8386	2
23101	2
19830	2
19010	2
22227	2
2457	2
9870	2
15791	2
16422	2
4970	2
16389	2
5921	2
9893	2
24114	2
15379	2
22896	2
12067	2
21696	2
12071	2
10118	2

```

16821    2
2719     2
17122    2
8912     2
20968    2
20458    2
19426    2
Name: label, Length: 4290, dtype: category
Categories (3, int64): [0, 1, 2]

```

```
In [41]: tf.confusion_matrix(df_train.label, predict_labels)
```

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/confusion_matrix.py:126:
Instructions for updating:
Use tf.cast instead.

```

```

W0513 08:52:16.111874 139779932305280 deprecation.py:323] From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/confusion_matrix.py:126:
Instructions for updating:
Use tf.cast instead.

```

```
Out[41]: <tf.Tensor 'confusion_matrix_1/SparseTensorDenseAdd:0' shape=(3, 3) dtype=int32>
```

```
In [0]: from sklearn.metrics import classification_report
report = classification_report(df_train.label, predict_labels)
```

```
In [43]: print(report)
```

	precision	recall	f1-score	support
0	0.83	0.87	0.85	1430
1	0.83	0.91	0.86	1430
2	0.97	0.82	0.89	1430
micro avg	0.87	0.87	0.87	4290
macro avg	0.87	0.87	0.87	4290
weighted avg	0.87	0.87	0.87	4290

```
In [45]: import seaborn
from sklearn.metrics import confusion_matrix
confusion_matrix = confusion_matrix(df_train.label, predict_labels)
matrix_proportions = np.zeros((3,3))
for i in range(0,3):
    matrix_proportions[i,:] = confusion_matrix[i,:]/float(confusion_matrix[i,:].sum())
names=['Hate','Offensive','Neither']
confusion_df = pd.DataFrame(matrix_proportions, index=names,columns=names)
```

```
plt.figure(figsize=(5,5))
seaborn.heatmap(confusion_df,annot=True,annot_kws={"size": 12},cmap='gist_gray_r',cbar=
plt.ylabel(r'True categories',fontsize=14)
plt.xlabel(r'Predicted categories',fontsize=14)
plt.tick_params(labelsize=12)
```

